

BOW Models

Vsevolod Dyomkin

prj-nlp

2020-04-02

Contents

- * BoW Concept
- * Generative & Discriminative Models
- * AP & NB Learning Algorithms
- * Feature Collection, Stopwords
- * A Number of Practical Examples
- * Simple Text Similarity Measures
- * TF-IDF

NLP Viewpoints

bag-of-words



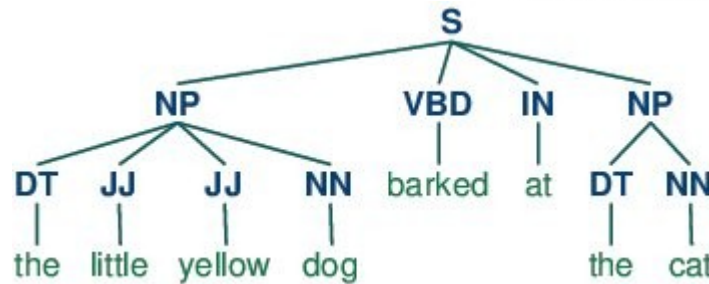
lexics

sequence

worse words warse wards wans wynds weans
weeds weens weens weens weens weens weens

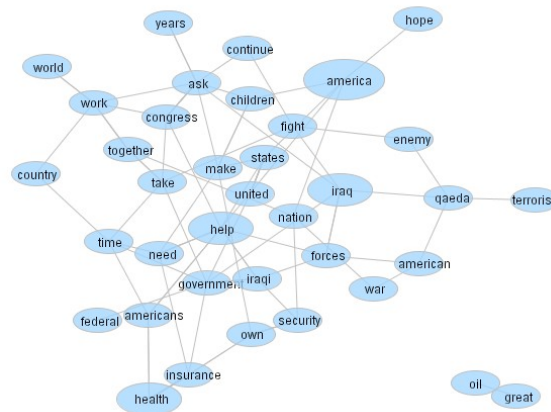
discourse

tree



syntax

graph






semantics

The Glorious BoW

- * Simplest model
- * Feature vector in N-dim space -
vector of words (with or w/o counts)
(N = dictionary size) -
a.k.a 1-hot representation

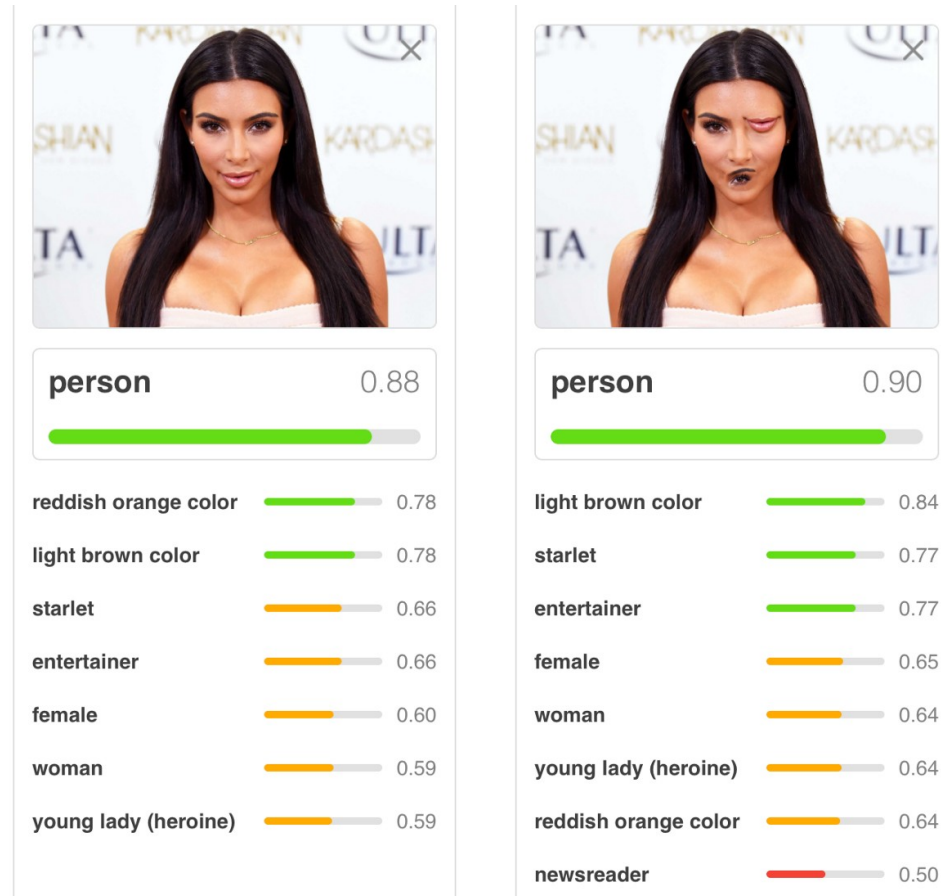
ONE-HOT ENCODING

	bread	yogurt	muffins
	1	0	0
	0	1	0
	0	0	1

365° DataScience


- * Position information disregarded
- * Works mostly for c12n

... not only for text



<https://hackernoon.com/capsule-networks-are-shaking-up-ai-heres-how-to-use-them-c233a0971952>

Generative vs Discriminative ML Models



Generative Models

- * Model joint probability of a sample and label:
 - can be used both to classify and generate
- * Introduce some structure (constraints)
- * That's why accuracy is usually asymptotically lower (but learn faster)
- * Examples:
 - Naive Bayes
 - GMM
 - HMM
 - PCFG
 - GAN

Generative vs Discriminative Models

<https://stats.stackexchange.com/questions/12421/generative-vs-discriminative>

- a) The generative model does indeed have a higher asymptotic error (as the number of training examples become large) than the discriminative model, but
- b) The generative model may also approach its asymptotic error much faster than the discriminative model — possibly with a number of training examples that is only logarithmic, rather than linear, in the number of parameters

<http://ai.stanford.edu/~ang/papers/nips01-discriminativegenerative.pdf>

Spam Identification

A 2-class whole text classification problem with a bias towards minimizing FPs.

Default approach - Rule-based (SpamAssassin)

Problems:

- scales poorly
- hard to reach arbitrary precision
- hard to rank the importance of complex features
- hard to interpret score and use it in upstream calculations



Apache SpamAssassin

“A Plan for Spam”

Proposed by Paul Graham


(<http://www.paulgraham.com/spam.html>)

1. Use the BoW approach
2. Use the Naive Bayes learning algorithm
3. Train on a balanced corpus

Initial results: Rec: 92%, Prec: 98.84%

Improved results: Rec: 99.5%, Prec: 99.97%

Bayes Rule



STATISTICALLY SPEAKING, IF YOU PICK UP A SEASHELL AND DON'T HOLD IT TO YOUR EAR, YOU CAN PROBABLY HEAR THE OCEAN.

Naive Bayes Classifier

$$P(Y|X) = P(Y) * P(X|Y) / P(X)$$

select $Y = \operatorname{argmax} P(Y|x)$

Naive step:

$$P(Y|x) = P(Y) * \prod_{\text{for all } x \text{ in } X} P(x|Y)$$

($P(x)$ is marginalized out because it's the same for all Y)

NB Model for Spam

madam	0.99	perl	0.01
promotion	0.99	python	0.01
republic	0.99	tcl	0.01
shortest	0.047225013	scripting	0.01
mandatory	0.047225013	morris	0.01
standardization	0.07347802	graham	0.01491078
sorry	0.08221981	guarantee	0.9762507
supported	0.09019077	cgi	0.9734398
people's	0.09019077	paul	0.027040077
enter	0.9075001	quite	0.030676773
quality	0.8921298	pop3	0.042199217
organization	0.12454646	various	0.06080265
investment	0.8568143	prices	0.9359873
very	0.14758544	managed	0.06451222
Valuable	0.82347786	difficult	0.071706355

<https://alexn.org/blog/2012/02/09/howto-build-naive-bayes-classifier.html>

The Value of Pre/Post-Processing

“Clever tricks”:

- title is more important than text
- text in the beginning is more important than at the end
- UNKs handling (spammers are smart)

Pre-processing:

- numbers pre-processing
- take only 15 most “interesting” words

...also: non-NLP features

NB Model for Lang ID

- * The problem of using words
- * Character ngrams to the rescue
- * Combining them

Sentiment Analysis

A 3-class whole-text¹ classification problem.

Default approach - Lexicon-based

Possible problems:

- ???

BoW Models for Sentiment

Features: words, bigrams


Models:

- * Multinomial NB
- * SVM with 2nd-order polynomial kernel
- * NBSVM

<https://www.aclweb.org/anthology/P12-2018>

BoW Fail Cases

- * polysemy
- * negation
- * neutralization
- * multiple sentiments
- * multiple objects
- * ambiguity
- * noise (errors)



Negation Examples

- * Morphological:

The food was **no** good.

I did **not** like them.

Their food was **without** any taste.

They **lack** good manners.

- * Syntactic:

If only their prices weren't that high!

I wish the food they served was more delicious.

Unlike The X, The Y has great service.

If they weren't rude, they wouldn't have lost their customers.

They are unlikely to improve.

False Negation

High prices were **no** surprise.

There is **no** reason to not like them.

It will bring us **nowhere**, but to success.

There's **no** doubt they are going to win the market.

The restaurant was **not** only cozy, but also located in a wonderful place.

Not only were the waiters rude, but they also brought the wrong dishes.

Neutralization

- * Morphological:

The X was **once** described as a leader in sales.

Earlier, The X used to put off the customers a lot.

- * Syntactic:

If they engage more customers, they will earn more.

All the hotels, **excluding** The X Hotel, were sued.

The restaurant was **neither good, nor bad**.

Multiple Sentiments

My sons loved The Playground. They are great, not like The Sandbox with their unsanitary kitchen. High prices were no surprise, though.

Ambiguity

The company is worth the words that were said earlier.

It tastes like beer.

It's in the same league with The Happiness Project, trust me.

Obama was right about it.

More BoW “Tricks”

- * normalization of special tokens
- * lemmatization/stemming
- * stopwords removal
- * filtering by “relevance”
(e.g. TF-IDF)
- * filtering by LM, parse, SRL...
- * combining words (negation, prepositions, NER...)

Stopwords



thelousylinguist
@lousylinguist



NLPers, stop removing stop words "just cuz". I repeated a text classification tutorial (analyticsvidhya.com/blog/2018/11/t...) but skipped the 'remove stop words' section and got a 2.8% INCREASE in accuracy. Stop words can improve your outcomes in many cases.

```
[ ] #again try to fit our model to see a big increase in accuracy.  
learn.fit_one_cycle(1, 1e-2)
```

☞ Total time: 00:36

epoch	train_loss	valid_loss	accuracy
1	0.547412	0.396379	0.896624

```
[37] #WITH STOP WORDS  
#again try to fit our model to see a big increase in accuracy.  
learn.fit_one_cycle(1, 1e-2)
```

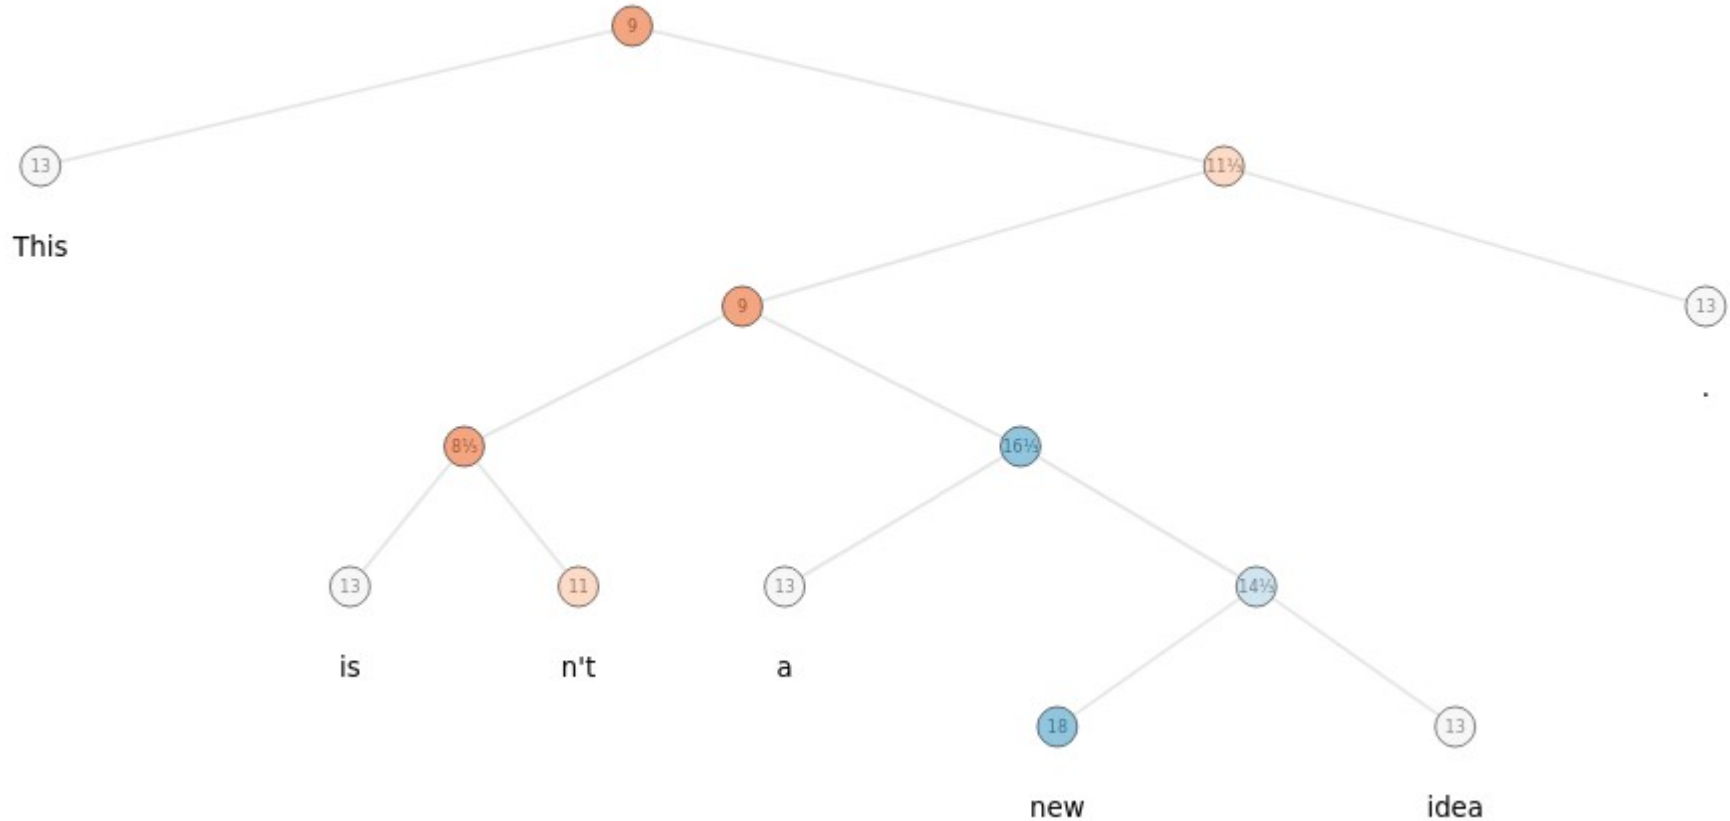
☞ Total time: 00:36

epoch	train_loss	valid_loss	accuracy
1	0.493484	0.285463	0.924051

1:30 AM · Nov 30, 2018 from Fairfield, CA · [Twitter for Android](#)

112 Retweets 445 Likes

Sentiment Treebank



<https://nlp.stanford.edu/sentiment/treebank.html>

Massive Multi-Class Problems

Problem: classification of user-generated product post according to a 1k labels catalog

Issues:

- class imbalance
- low representation of long-tail classes

Solutions:

- multi-level classification
- 1 model per level with the output classes limited to a particular subset
- AP learning algorithm

Discriminative Models

- * Model conditional probability of label, given a sample:
 - can be used only to classify
- * Training is direct
- * Examples:
 - kNN
 - Perceptron & Averaged Perceptron
 - Logistic Regression (aka MaxEnt)
 - AROW
 - SVM
 - CRF
 - Feed-forward Neural Nets (FNN)

(Averaged) Perceptron

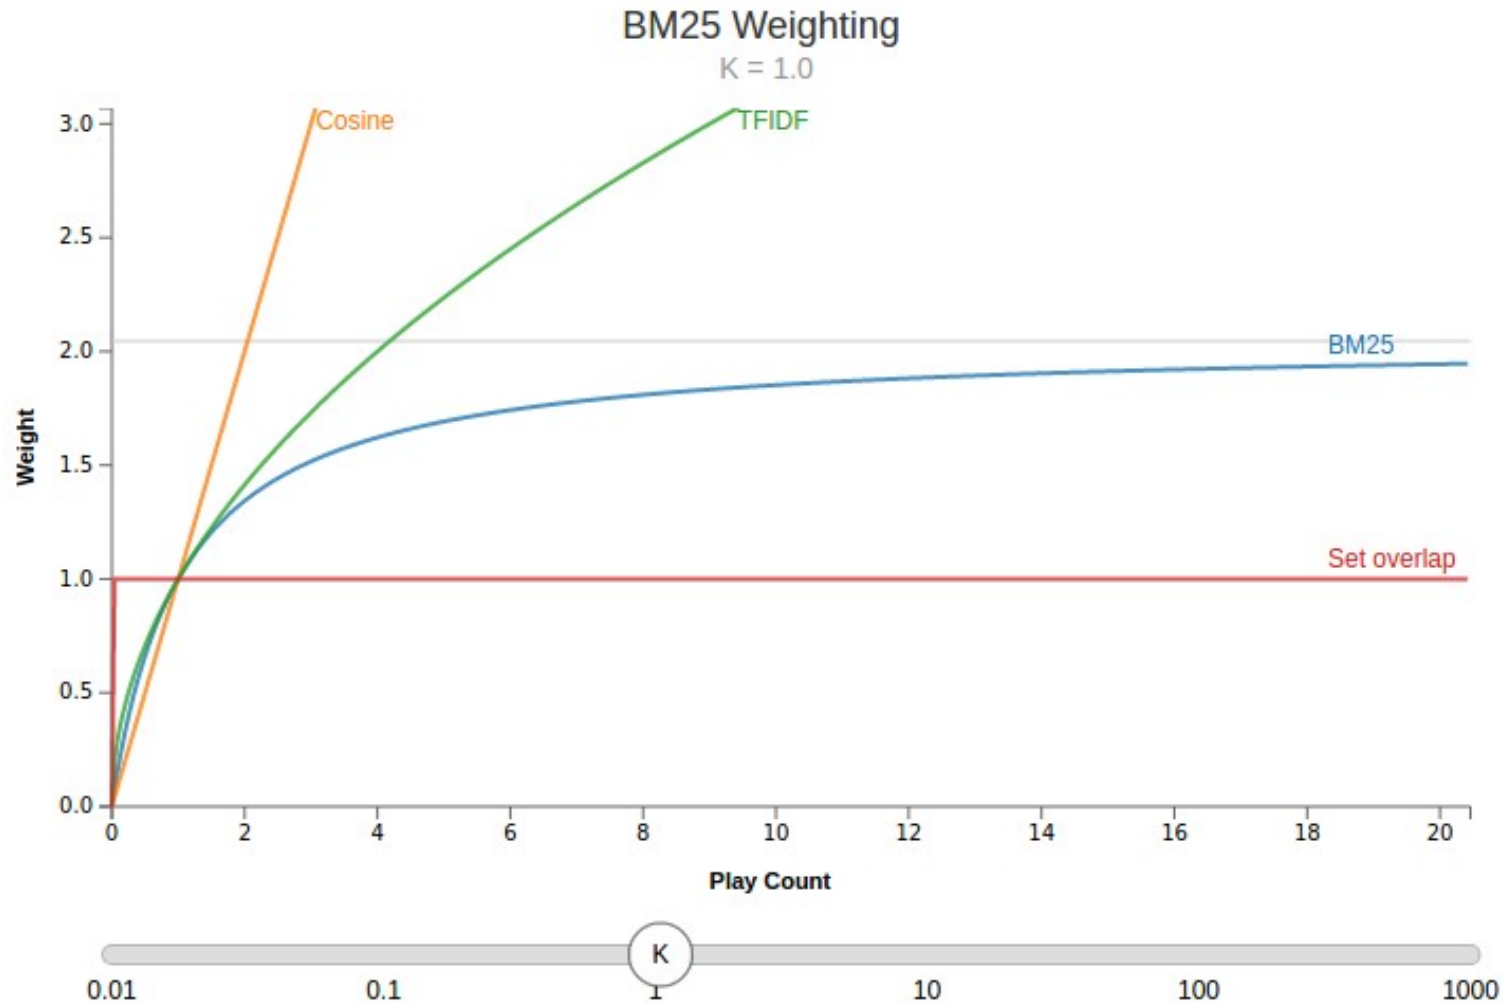
- * Simplest linear discriminative model
- * On-line learning
- * When averaged — ensemble, asymptotic optimality

Perceptron learning rule:

```
def train(self, nr_iter, examples):  
    for i in range(nr_iter):  
        for features, true_tag in examples:  
            guess = self.predict(features)  
            if guess != true_tag:  
                for f in features:  
                    self.weights[f][true_tag] += 1  
                    self.weights[f][guess] -= 1  
        random.shuffle(examples)
```

<https://explosion.ai/blog/part-of-speech-pos-tagger-in-python>

Similarity Metrics



<http://www.benfrederickson.com/distance-metrics/>

Jaccard Similarity

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}.$$

Cosine Similarity

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

BM25 Similarity

```
def bm25_tf_weight(item):  
    return item * (K1 + 1.0) / (K1 + item)
```

By changing the value of $K1$ in this function, we can change the shape to go between the step function used in the Jaccard distance ($K1 = 0$) and the linear weighting used in the Cosine distance ($K1 = +\infty$).

The other major change with BM25 is how length normalization is handled. Term counts are scaled by the ratio of the size of the document to the average size. But since sometimes it makes sense to prefer longer documents, BM25 also introduces a parameter B which controls how much influence the length normalization has on the results.

TF-IDF

A classic IR technic for ranking relevancy

$$w_{x,y} = tf_{x,y} \times \log \left(\frac{N}{df_x} \right)$$

TF-IDF

Term x within document y

$tf_{x,y}$ = frequency of x in y

df_x = number of documents containing x

N = total number of documents

TF-IDF

Variants of term frequency (TF) weight

weighting scheme	TF weight
binary	0, 1
raw count	$f_{t,d}$
term frequency	$f_{t,d} / \sum_{t' \in d} f_{t',d}$
log normalization	$1 + \log(f_{t,d})$
double normalization 0.5	$0.5 + 0.5 \cdot \frac{f_{t,d}}{\max_{\{t' \in d\}} f_{t',d}}$
double normalization K	$K + (1 - K) \frac{f_{t,d}}{\max_{\{t' \in d\}} f_{t',d}}$

TF-IDF

Variants of inverse document frequency (IDF) weight

weighting scheme	IDF weight ($n_t = \{d \in D : t \in d\} $)
unary	1
inverse document frequency	$\log \frac{N}{n_t} = -\log \frac{n_t}{N}$
inverse document frequency smooth	$\log \left(1 + \frac{N}{n_t} \right)$
inverse document frequency max	$\log \left(\frac{\max_{t' \in d} n_{t'}}{1 + n_t} \right)$
probabilistic inverse document frequency	$\log \frac{N - n_t}{n_t}$

Science Pulse



Keyphrase	Weight	Keyphrase	Weight
massive multiple input output	7,25	все буде добре	1
long short term memory architecture	5,12	коли тебе нема	1
Live action virtual reality games	3,15	небо над дніпром	1
low rank hankel matrix completion	3,04	хочу напиться тобою	0,78
multi point wireless energy transmission	3,01	жити без мети	0,78
tree augmented naive bayes classifier	2,89	мила моя сьюзі	0,78
long short term memorized fusion	2,15	тінь твого тіла	0,75
fine grained entity type classification	1,51	коли настане день	0,75
high speed railway communication systems	1,27	кожну хвилину життя	0,75
partially observable markov decision process	1,13	коли тобі важко	0,75

<https://aiukraine.com/wp-content/uploads/2016/09/Tetiana-Kodliuk.pdf>

RAKE

RAKE short for Rapid Automatic Keyword Extraction algorithm, is a domain independent keyword extraction algorithm which tries to determine key phrases in a body of text by analyzing the frequency of word appearance and its co-occurrence with other words in the text.

RAKE Pipeline:

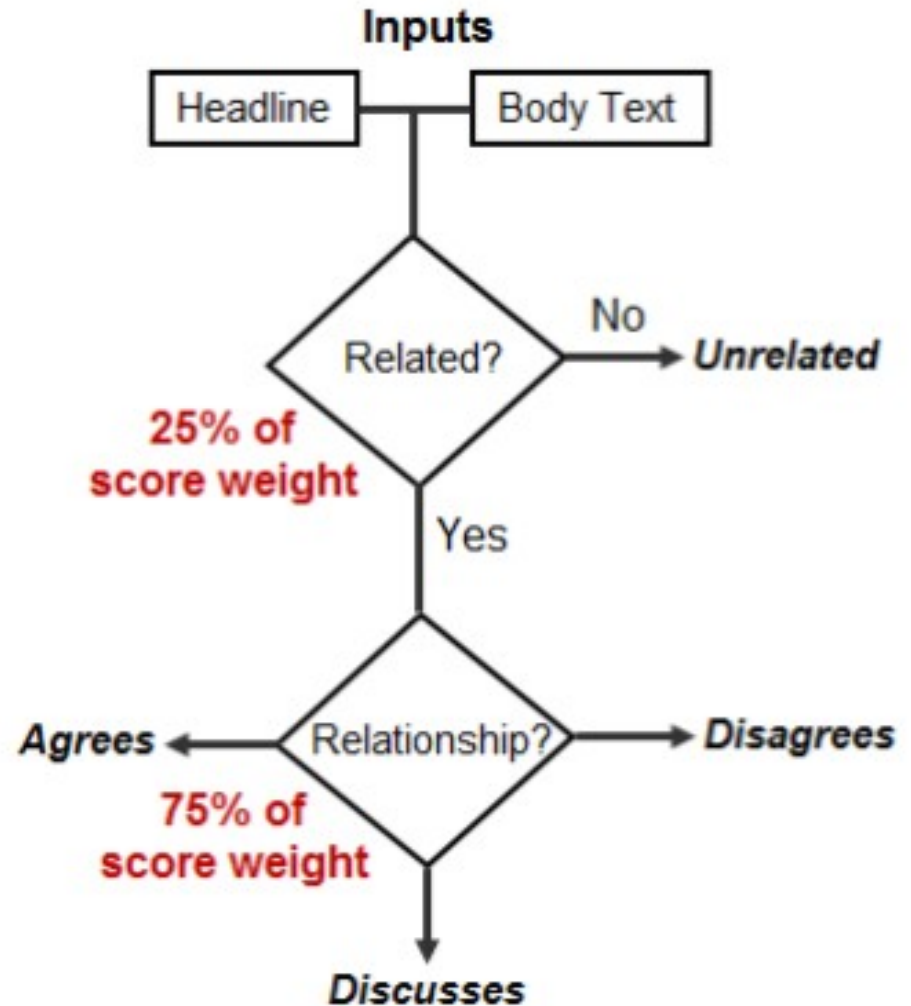
- * Partition the text by stop words and punctuation
- * Creates a co-occurrence matrix of terms.
- * For each content word,
count $\text{deg}(\text{word}) / \text{freq}(\text{word})$.
- * For each key phrase, sum scores of words.
- * Return the top 1/3 of key phrases.

<https://pypi.org/project/rake-nltk/>

Stance Detection

A 4-class whole text
Hierarchical
classification problem:

- * unrelated,
- * related:
 - discuss
 - agree
 - disagree



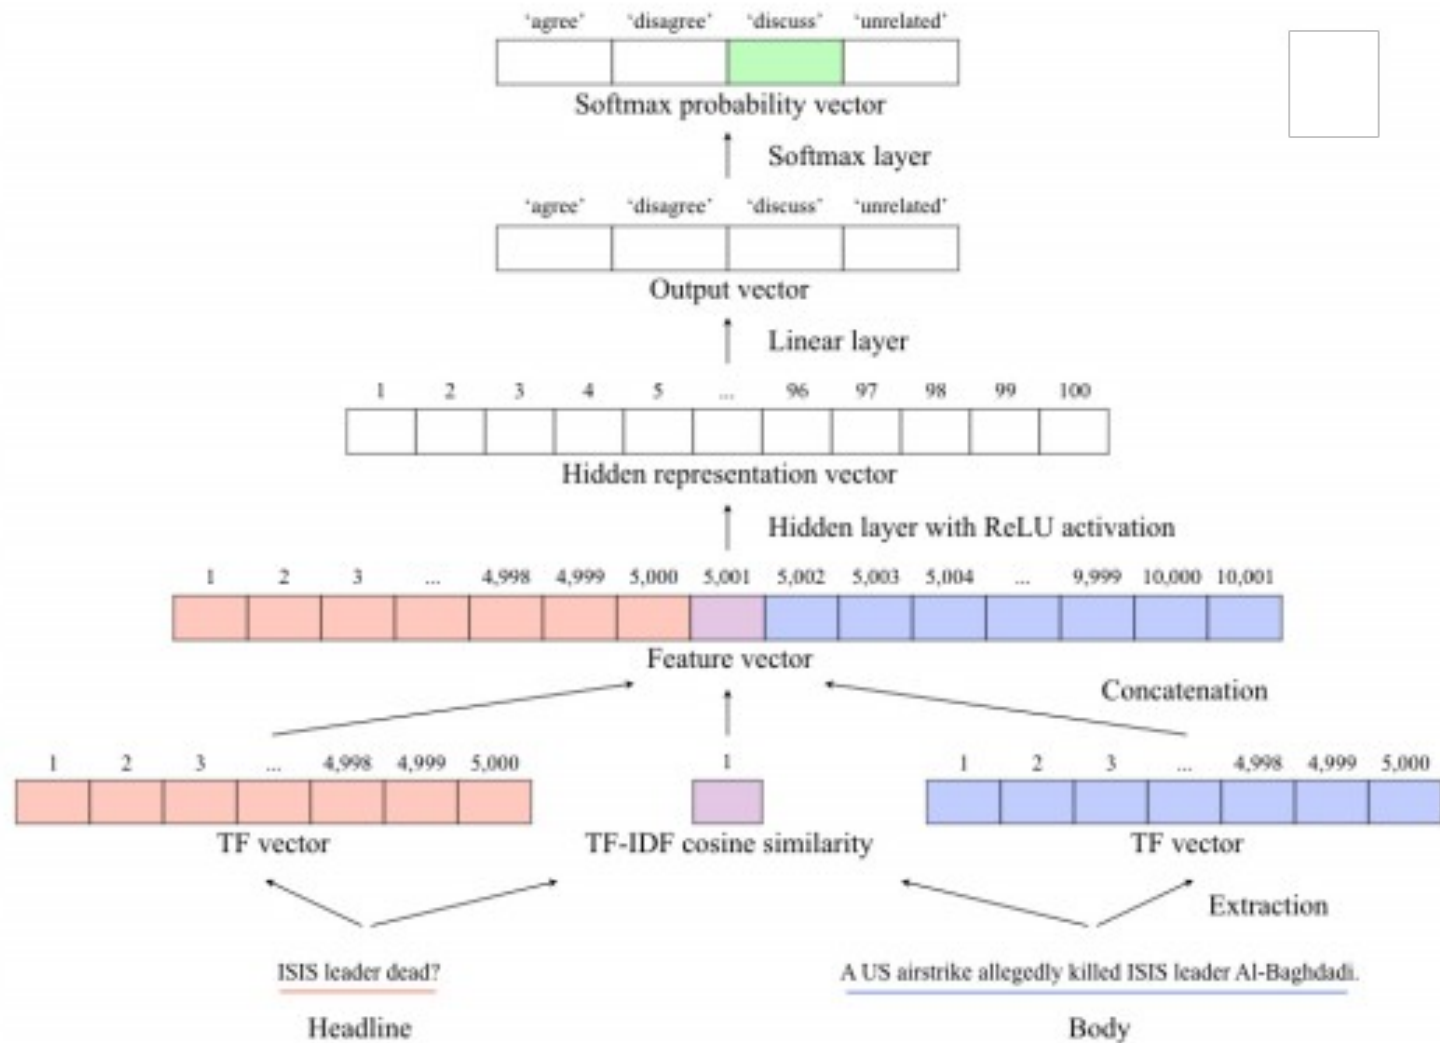
<https://github.com/FakeNewsChallenge/fnc-1>

TF-IDF Cosine Similarity Baseline

	agree	disagree	discuss	unrelated
agree	94	14	111	543
disagree	13	27	9	113
discuss	11	31	607	1151
unrelated	379	91	650	5778

Score: 2219.75 out of 4448.5 (49.898842306395416%)

BoW-MLP Model



BoW Recap

Pros:

- + simple
- + easy to compute
- + flexible
- + doesn't require lots of data
- + with lots of data works very well

Cons:

- doesn't capture order
- hard to capture inter-word relations
- hard to scale to real-world vocabularies
- poor generalization

Read More

Sentiment analysis:

- http://www.datasciencecentral.com/profiles/blogs/test?xg_source=activity
- <http://ataspinar.com/2015/11/16/text-classification-and-sentiment-analysis/>
- <http://ataspinar.com/2016/02/15/sentiment-analysis-with-the-naive-bayes-classifier/>
- <https://www.cs.uic.edu/~liub/FBS/Sentiment-Analysis-tutorial-AAAI-2011.pdf>

BoW in CV:

- http://cs.nyu.edu/~fergus/teaching/vision_2012/9_BoW.pdf

BoW on Extremely Small Datasets:

- <https://towardsdatascience.com/text-classification-with-extremely-small-datasets-333d322caee2>

Beyond Cosine Similarity

- <https://stefansavev.com/blog/beyond-cosine-similarity/>