

Neural Nets for NLP

Vsevolod Dyomkin
prj-nlp-2020


Limitations of “Classic” NLP

- * categorical (1-hot) features
- * extra large feature spaces
- * UNK problems
- * complicated feature engineering
- * difficult domain adaptation
- * need for markovization in sequence models
- * what else?

Neural Nets to the Rescue

RBM CNN
MLP DBN SOM
SNN RNN
Hopfield GAN
VAE Capsule

Terminology



Deep Learning?

<https://blog.keras.io/the-limitations-of-deep-learning.html>

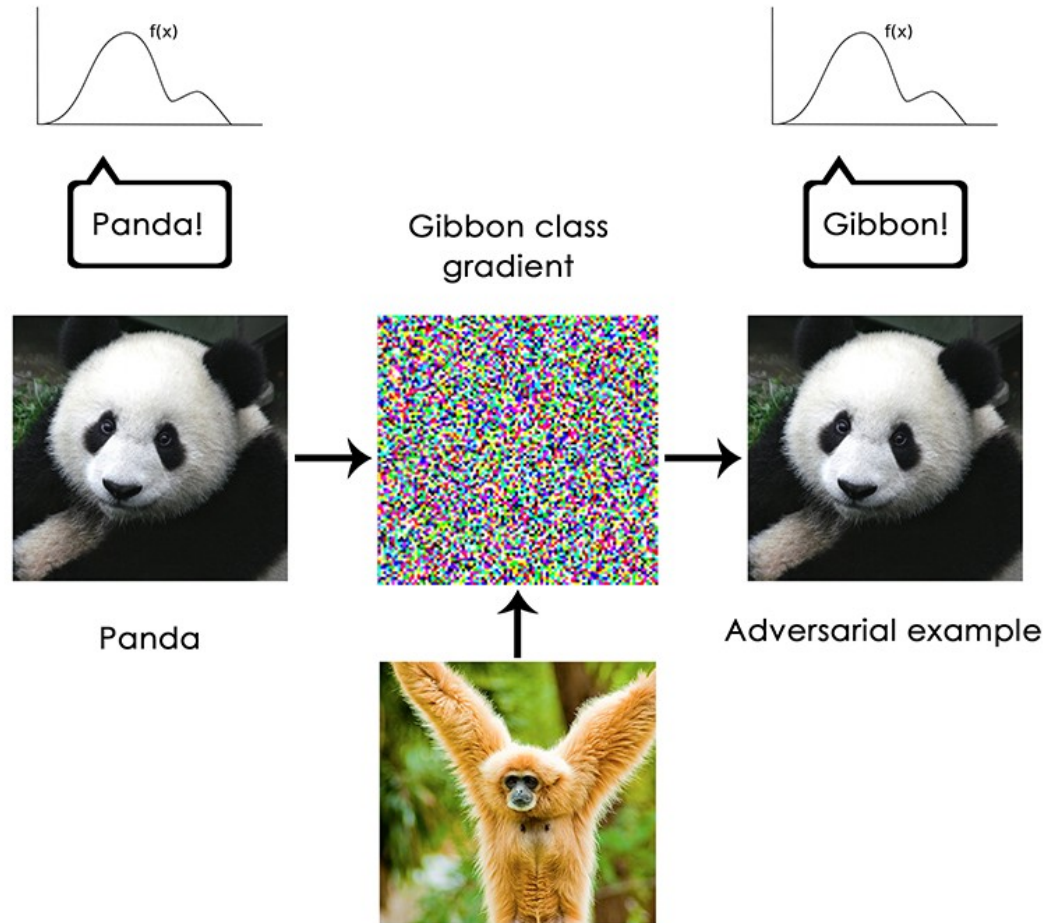
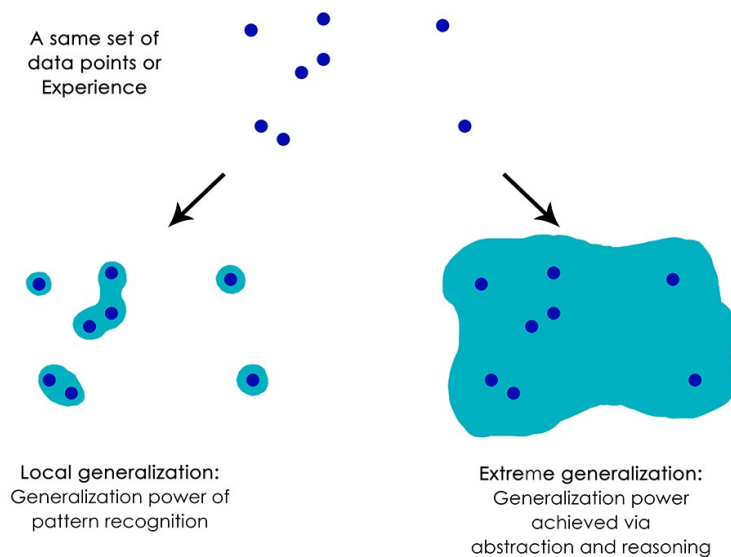
A Geometric View of DL

In deep learning, everything is a vector, i.e. everything is a point in a geometric space. Model inputs (it could be text, images, etc) and targets are first "vectorized", i.e. turned into some initial input vector space and target vector space. Each layer in a deep learning model operates one simple geometric transformation on the data that goes through it. Together, the chain of layers of the model forms one very complex geometric transformation, broken down into a series of simple ones. This complex transformation attempts to map the input space into the target space, one point at a time. This transformation is parameterized by the weights of the layers, which are iteratively updated based on how well the model is currently performing. A key characteristic of this geometric transformation is that it must be differentiable, which is required in order for us to be able to learn its parameters via gradient descent. Intuitively, this means that the geometric morphing from inputs to outputs must be smooth and continuous—a significant constraint.


The Limitations of DL




The boy is holding a baseball bat.




Types of Neural Networks



Single neuron: perceptron,
linear / logistic regression



Feed-forward network
(no cycles) -- non-linear
classification & regression



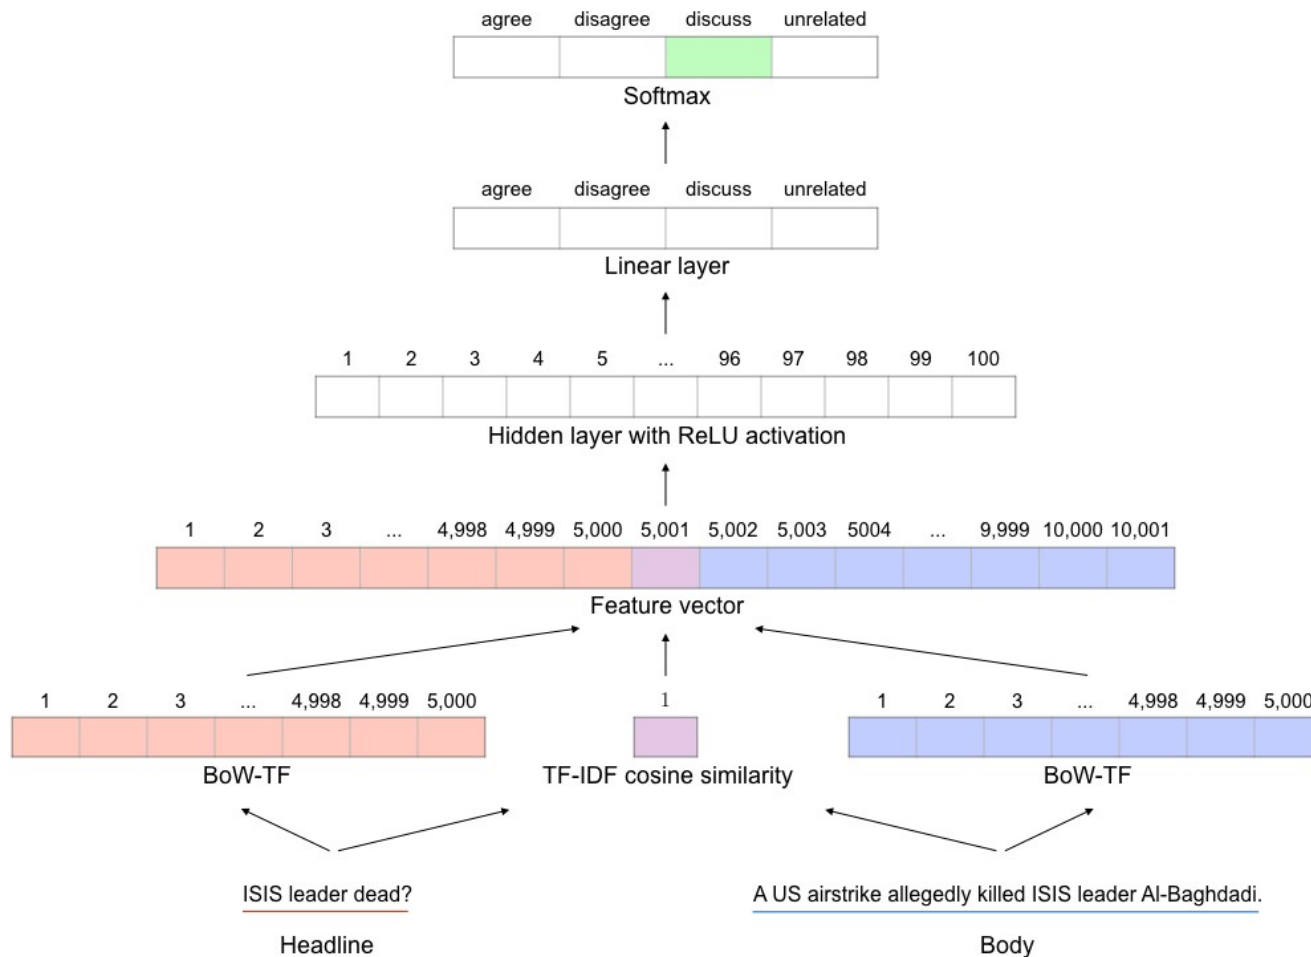
Symmetric (RBM)
unsupervised, trained
to maximize likelihood
of input data
a mixture model

~~MLPs~~ FNNs

- * computational graph
- * composed of various layers
- * backprop for learning
- * GD for optimization



Example: FNC-1




<https://github.com/uc1mr/fakenewschallenge>

Layers

- * input
- * fully-connected
- * convolutional
- * non-linearities
- * regularization
- * output


Input Layers

- * 1-hot
- * embedding
- * mixed





Fully-Connected Layers

* linear transformation




$$\text{output} = f(\text{Weights} \times \text{input} + \text{bias})$$



Convolutional Layers



* apply mask

+ pooling (max, mean,...)




Nonlinearities

- * sigmoid/logistic
- * tanh
- * ReLU/SeLU/ELU/leakyReLU/...
- * softplus
- * swish



- * maxout




SIGMOID / LOGISTIC

ADVANTAGES

- **Smooth gradient**, preventing “jumps” in output values.
- **Output values bound** between 0 and 1, normalizing the output of each neuron.
- **Clear predictions**—For X above 2 or below -2, tends to bring the Y value (the prediction) to the edge of the curve, very close to 1 or 0. This enables clear predictions.

DISADVANTAGES

- **Vanishing gradient**—for very high or very low values of X , there is almost no change to the prediction, causing a vanishing gradient problem. This can result in the network refusing to learn further, or being too slow to reach an accurate prediction.
- **Outputs not zero centered.**
- **Computationally expensive**



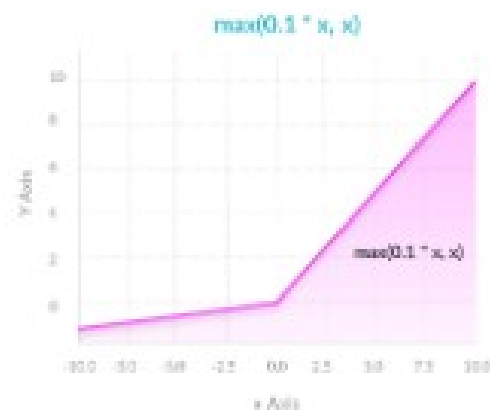
RELU (RECTIFIED LINEAR UNIT)

ADVANTAGES

- **Computationally efficient**—allows the network to converge very quickly
- **Non-linear**—although it looks like a linear function, ReLU has a derivative function and allows for backpropagation

DISADVANTAGES

- **The Dying ReLU problem**—when inputs approach zero, or are negative, the gradient of the function becomes zero, the network cannot perform backpropagation and cannot learn.




LEAKY RELU

ADVANTAGES

- **Prevents dying ReLU problem**—this variation of ReLU has a small positive slope in the negative area, so it does enable backpropagation, even for negative input values
- Otherwise like ReLU

DISADVANTAGES

- **Results not consistent**—leaky ReLU does not provide consistent predictions for negative input values.




TANH / HYPERBOLIC TANGENT

ADVANTAGES

- **Zero centered**—making it easier to model inputs that have strongly negative, neutral, and strongly positive values.
- Otherwise like the Sigmoid function.

DISADVANTAGES

- Like the Sigmoid function




SWISH


Swish is a new, self-gated activation function discovered by researchers at Google. According to their [paper](#), it performs better than ReLU with a similar level of computational efficiency. In experiments on ImageNet with identical models running ReLU and Swish, the new function achieved top -1 classification accuracy 0.6-0.9% higher.

Regularization Layers


- * nonlinearity regularization
- * dropout
- * dropconnect



No-Drop Network



DropOut Network




DropConnect Network

Output Layers

* softmax/hierarchical softmax

Multi-Class Classification with NN and SoftMax Function



Loss Functions

- * maximum likelihood estimation (MLE) — Cross Entropy (Log loss)

$$CE = - \sum_i^C t_i \log(s_i)$$

- * max margin objective - Hinge loss

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \max(0, m - y^{(i)} \cdot \hat{y}^{(i)})$$

- * KL-divergence loss

$$\begin{aligned} \mathcal{L} &= \frac{1}{n} \sum_{i=1}^n \mathcal{D}_{KL}(y^{(i)} || \hat{y}^{(i)}) \\ &= \frac{1}{n} \sum_{i=1}^n \left[y^{(i)} \cdot \log \left(\frac{y^{(i)}}{\hat{y}^{(i)}} \right) \right] \\ &= \underbrace{\frac{1}{n} \sum_{i=1}^n (y^{(i)} \cdot \log(y^{(i)}))}_{\text{entropy}} - \underbrace{\frac{1}{n} \sum_{i=1}^n (y^{(i)} \cdot \log(\hat{y}^{(i)}))}_{\text{cross-entropy}} \end{aligned}$$

https://isaacchanghau.github.io/post/loss_functions/

Backprop

- * efficient way to compute derivatives (using DP)
- * automatic & symbolic differentiation

<https://colah.github.io/posts/2015-08-Backprop/>

Optimization Algorithm

- * gradient descent
- * SGD (+minibatch)
- * Momentum
- * Adagrad/Adadelata/...
- * Adam

<https://ruder.io/optimizing-gradient-descent/>

Example: FNN for Adjective Ordering

Input: noun & 2 adjectives
(embeddings)

Hidden: 4 ReLU FC-layers

Output: Sigmoid (Softmax)

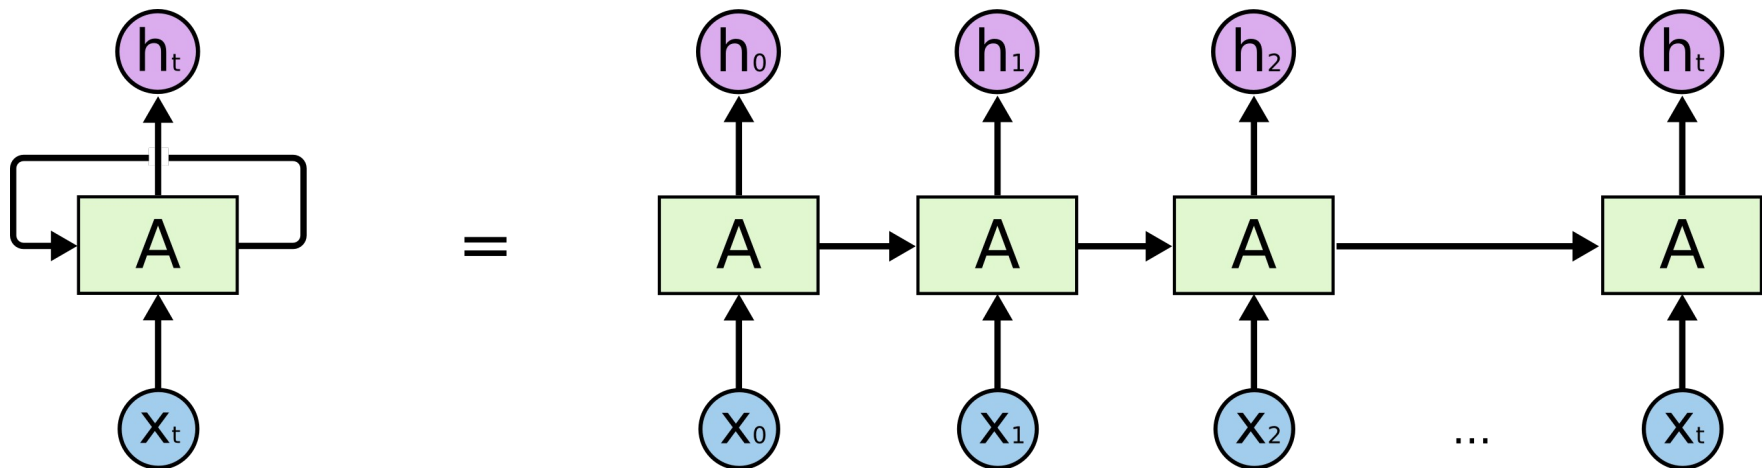
Run with both variants of order
and select the better score

FNNs Recap

- + nonlinear, flexible
- + efficient training
- + allows to use embeddings
- fixed input
- limited context

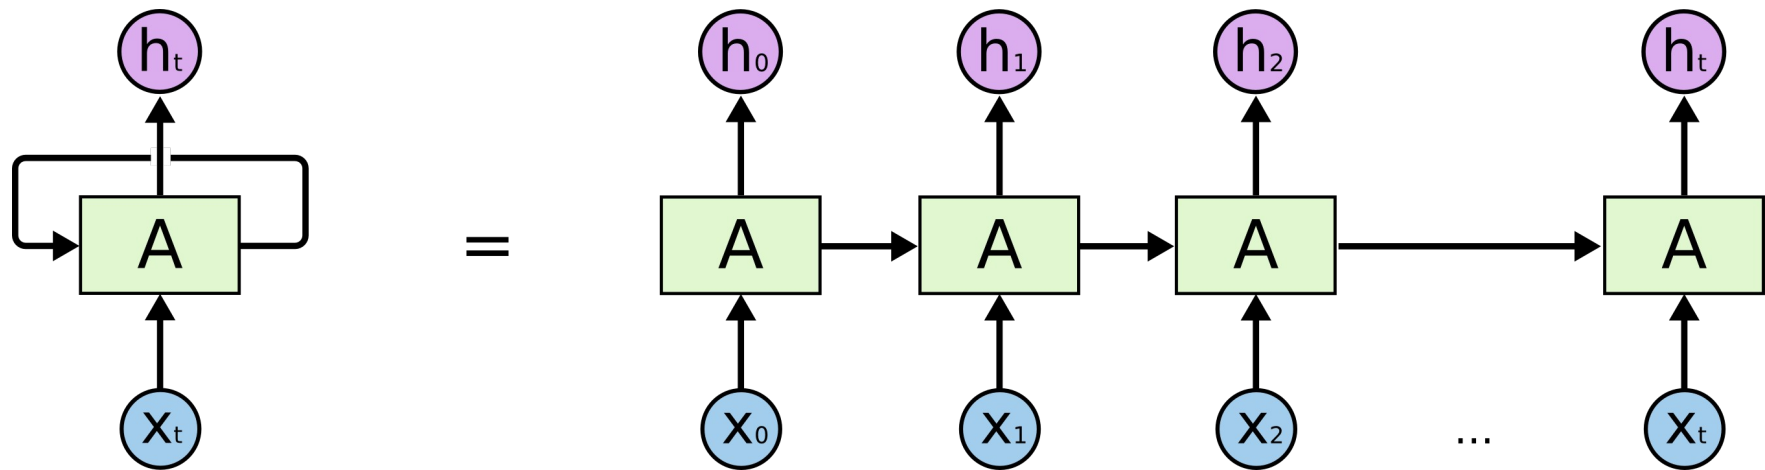
RNNs to the Rescue

- * add previous state to input
- * backpropagate through time (truncated - BPTT)



RNNs to the Rescue


- * add previous state to input
- * backpropagate through time



- * not so easy:
 - vanishing gradients
 - exploding gradients


LSTM

specifically designed to
remember long-term dependencies




<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

LSTM Cell State




LSTM Forget Gate




$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

LSTM Remember Gate



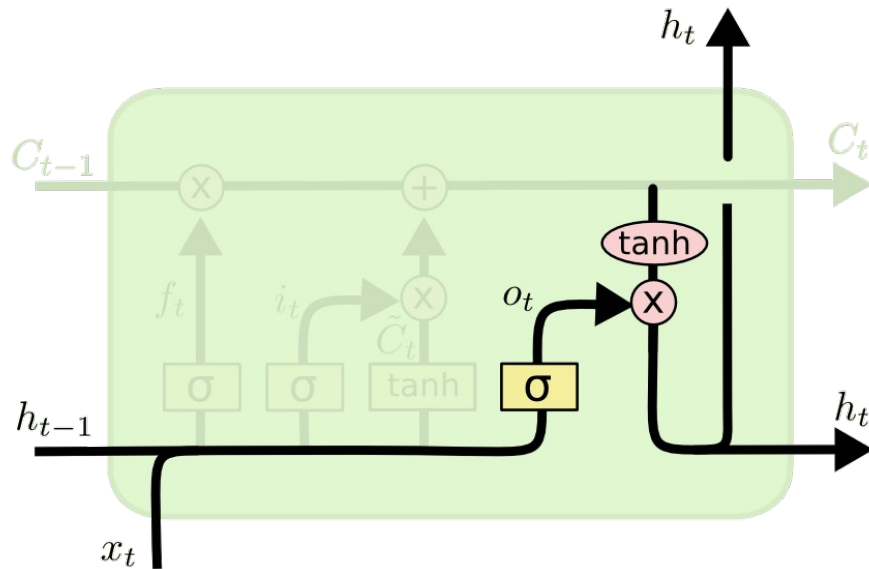
$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

LSTM State Update



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$


LSTM Output



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

LSTM Peephole Connections




$$f_t = \sigma (W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma (W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i)$$

$$o_t = \sigma (W_o \cdot [C_t, h_{t-1}, x_t] + b_o)$$

Gated Recurrent Unit (GRU)



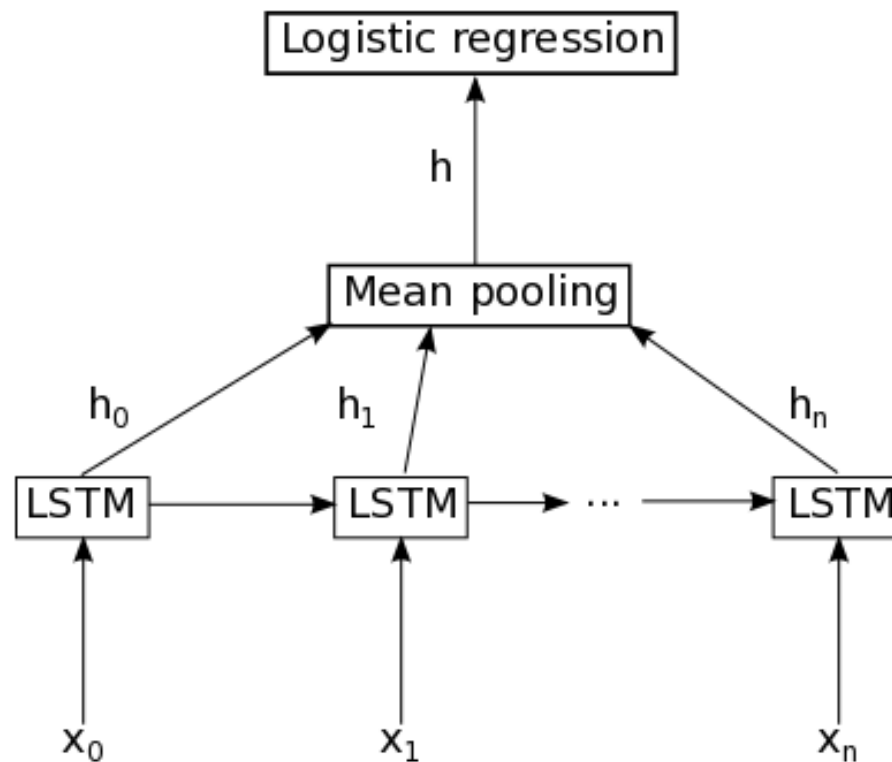
$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

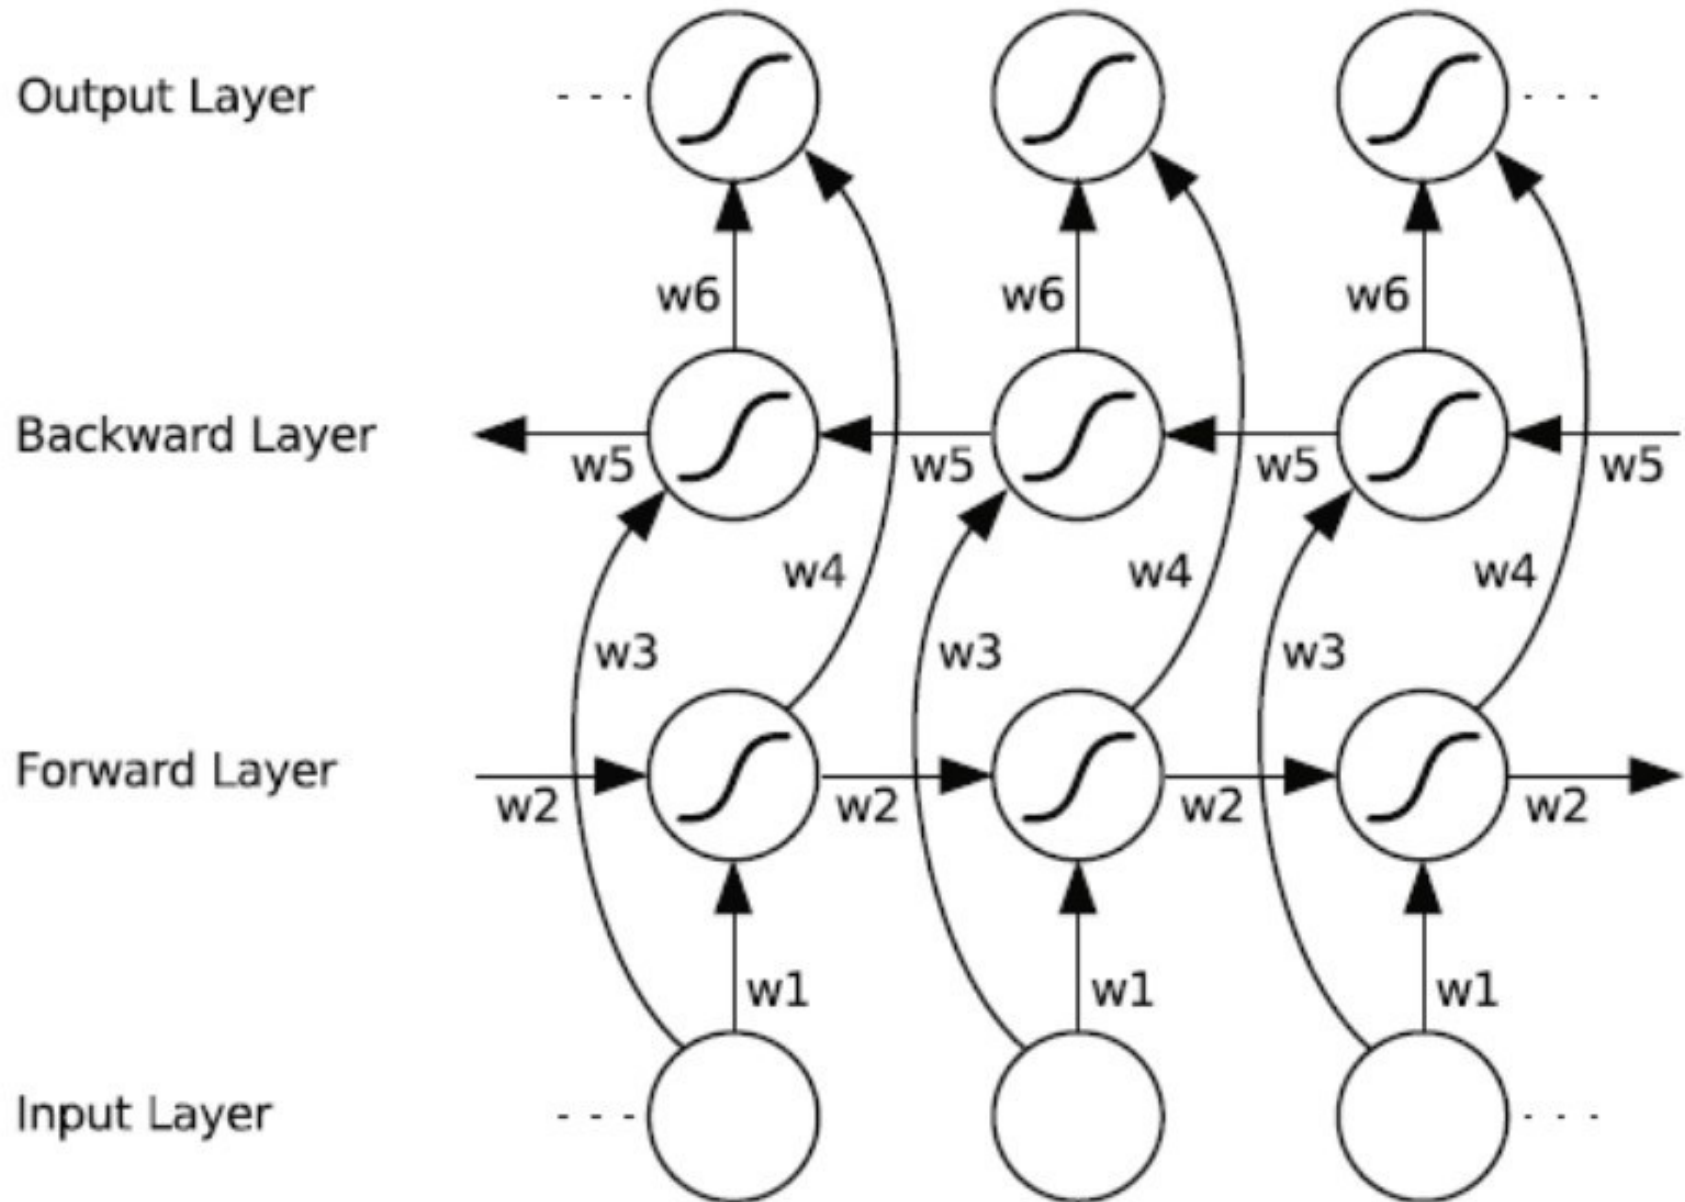
$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

LSTM Example: Sentiment Analysis

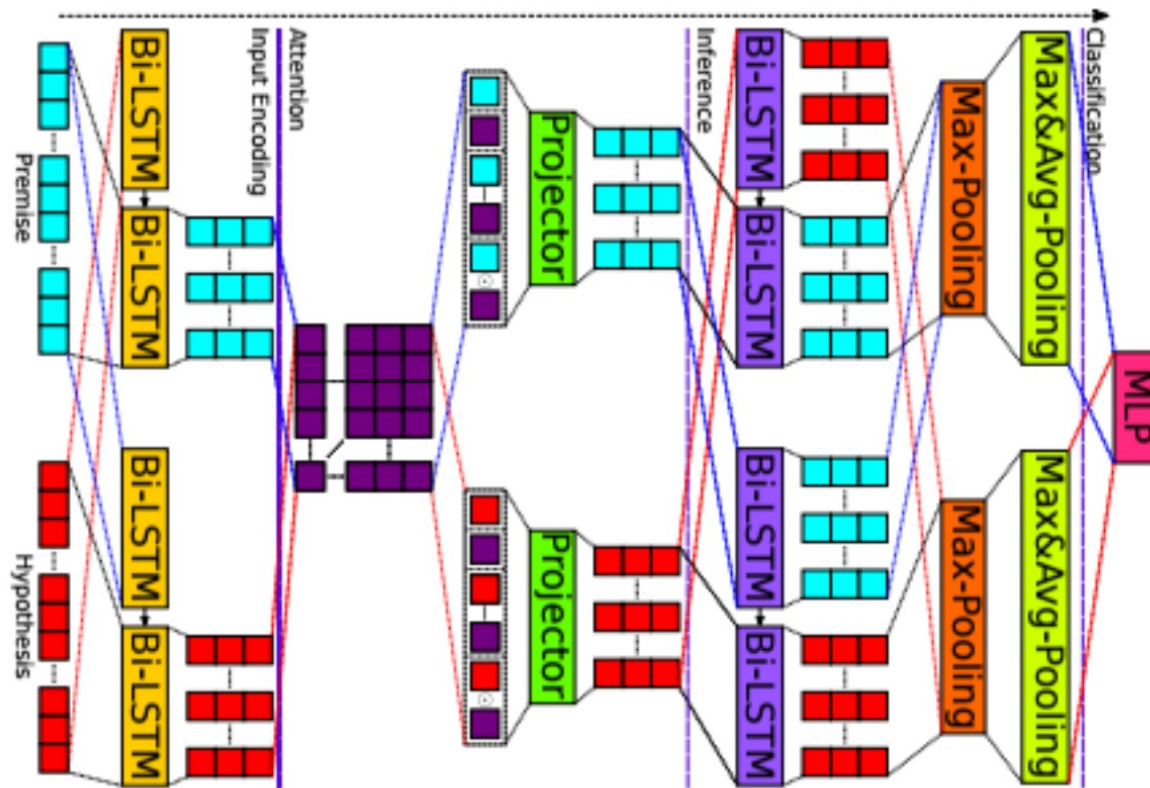


<http://deeplearning.net/tutorial/lstm.html>

BiLSTM



BiLSTM Example: Machine Reading

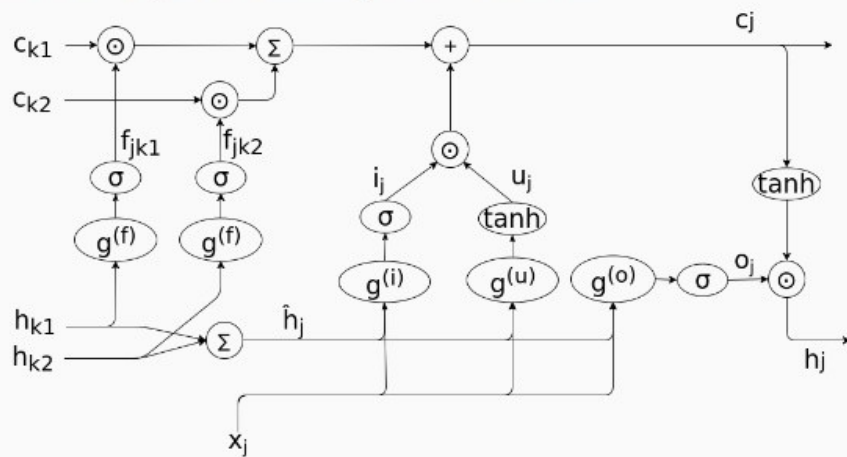


<https://arxiv.org/pdf/1802.05577.pdf>

TreeLSTM

Child-sum tree LSTM

Children outputs and memory cells are summed

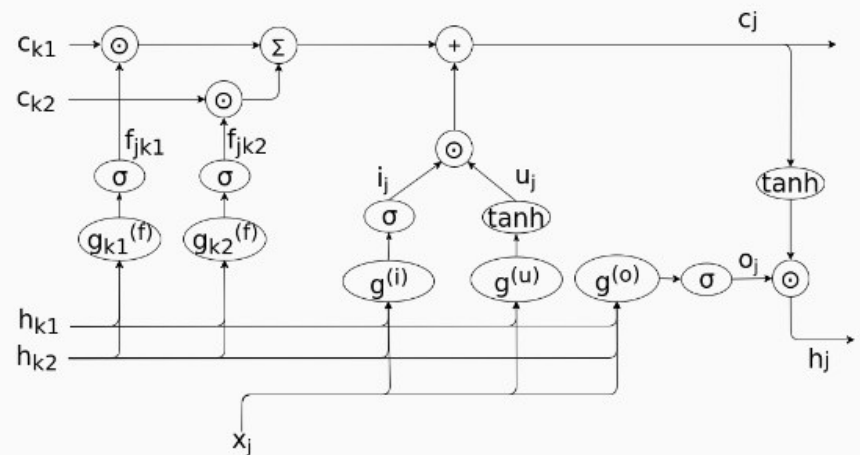


Child-sum tree LSTM at node j with children k_1 and k_2

10

N-ary tree LSTM

Given $g_k^{(n)}(x_t, h_{l_1}, \dots, h_{l_N}) = W^{(n)}x_t + \sum_{l=1}^N U_{kl}^{(n)}h_{l_l} + b^{(n)}$



Binary tree LSTM at node j with children k_1 and k_2

12

<https://www.slideshare.net/tuvistavie/tree-lstm>

TreeLSTM Example

Semantic relatedness

Task

Predict similarity score in $[1, K]$ between two sentences

Method

Similarity between sentences L and R annotated with score $\in [1, 5]$

- Produce representations h_L and h_R
- Compute distance h_+ and angle h_\times between h_L and h_R
- Compute score using fully connected NN

$$h_s = \sigma \left(W^{(\times)} h_\times + W^{(+)} h_+ + b^{(h)} \right)$$

$$\hat{p}_\theta = \text{softmax} \left(W^{(p)} h_s + b^{(p)} \right)$$

$$\hat{y} = r^T \hat{p}_\theta$$

$$r = [1, 2, 3, 4, 5]$$

- Error is computed using KL-divergence

LSTM Deficiencies

- * computation not parallelizable
- * neuron interpretability

Improvements/alternatives:

- * RAN (Recurrent Additive Network)
<https://arxiv.org/pdf/1705.07393.pdf>
- * Janet (just the forget gate)
<https://arxiv.org/pdf/1804.04849.pdf>
- * QRNN (Quasi-recurrent Neural Network)
<https://goo.gl/NUx7VC>

BiLSTM+attention as SOTA

“Basically, if you want to do an NLP task, no matter what it is, what you should do is throw your data into a Bi-directional long-short term memory network, and augment its information flow with the attention mechanism.”

— Chris Manning

<https://twitter.com/mayurbhangale/status/988332845708886016>


Convolutional Neural Networks (CNNs)

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

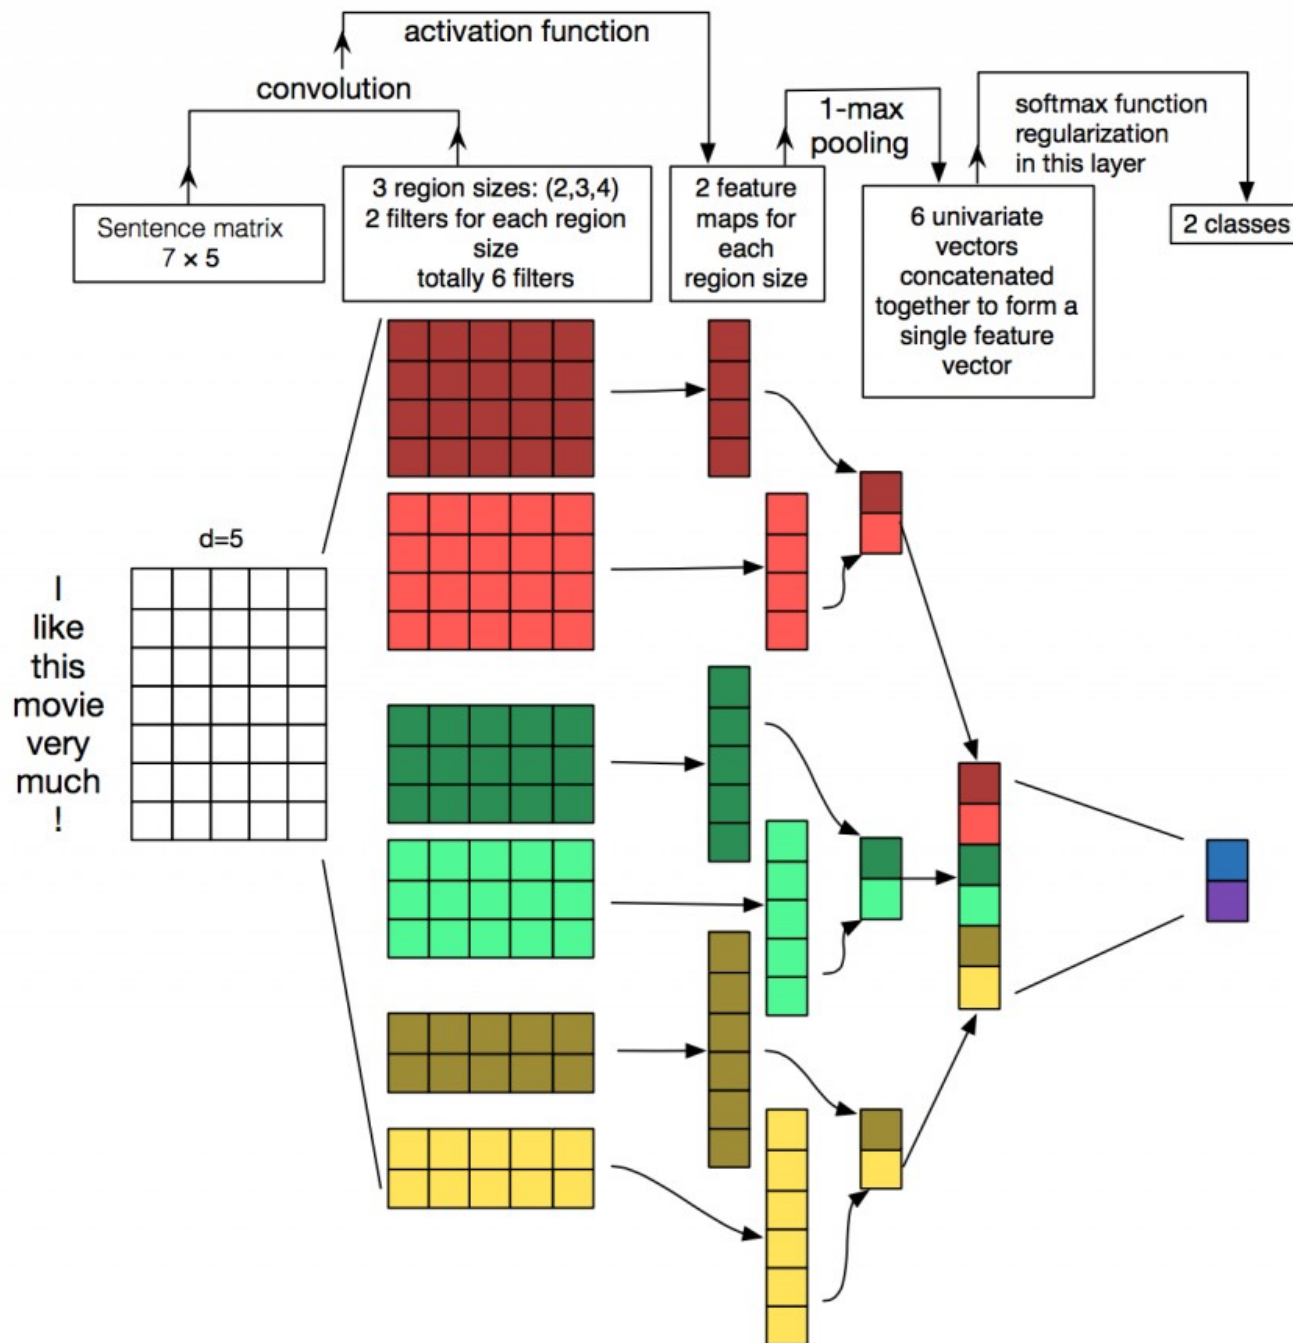
Image

4		

Convolved
Feature




- * convolutions can do the same as RNNs but faster
- * any part of a sentence can influence the semantics of a word
So we want our network to see the entire input at once
- * getting that big a receptive can make gradients vanish and our networks fail
- * we can solve the vanishing gradient problem with DenseNets or Dilated Convolutions
- * use “deconvolutions” to generate arbitrarily long outputs




CNN Parameters

Wide vs narrow convolutions:



Strides:




Pooling:

- max
- mean

CNN Example

Classification (sentiment et al.)

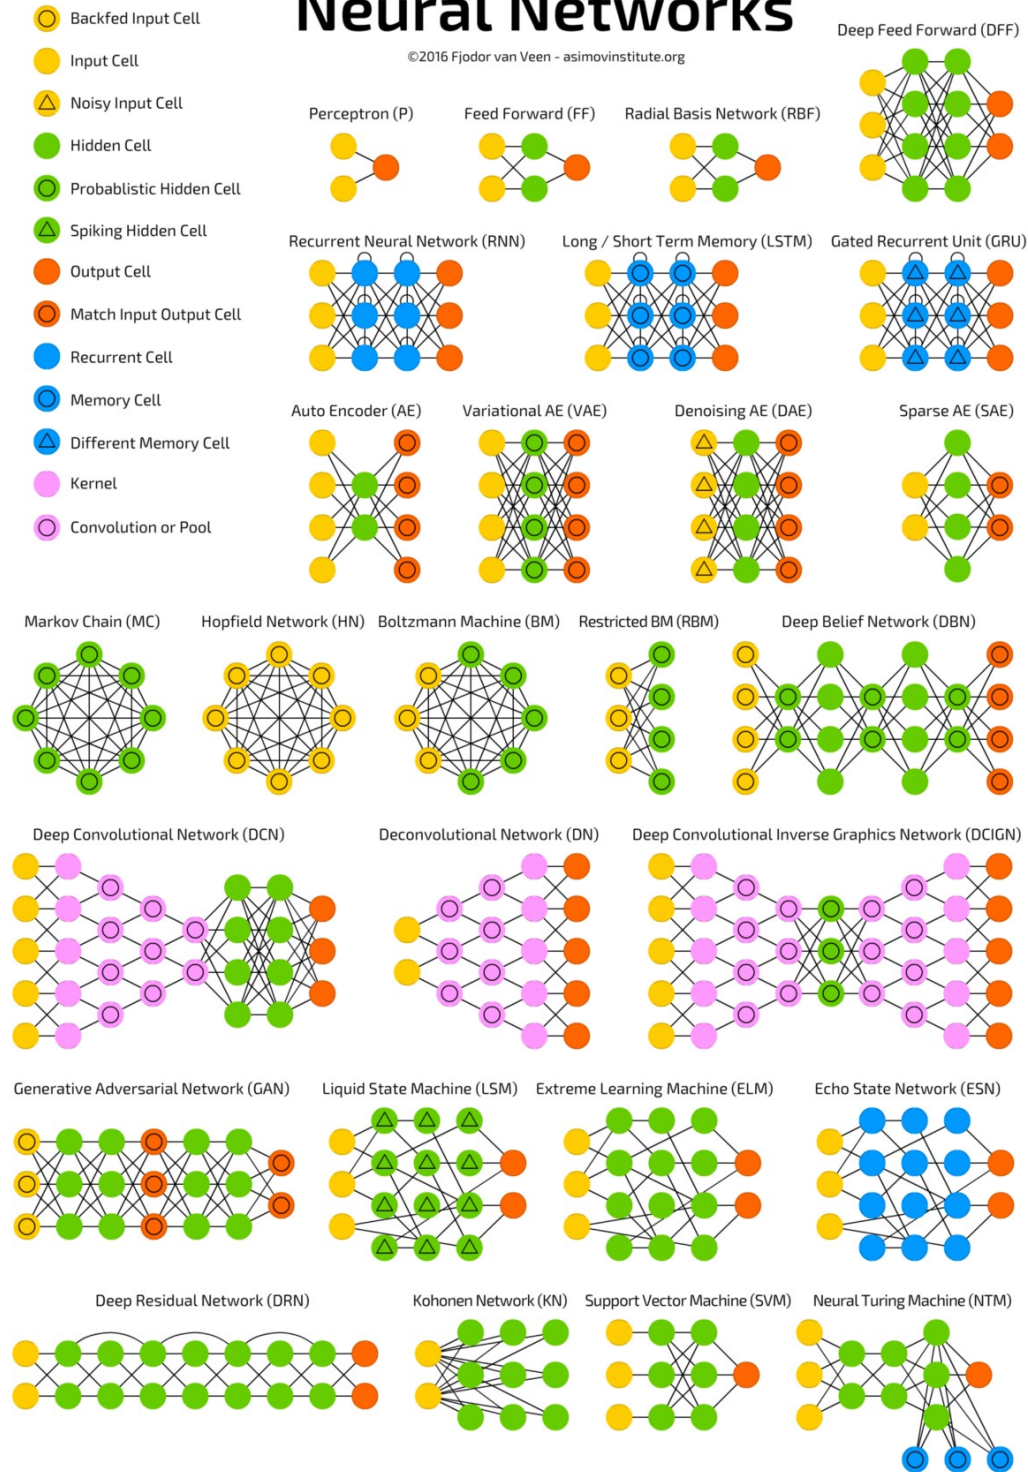


CNNs Pros & Cons

- + fast (and furious)
- + don't forget
- + view the whole input at once
- + can reuse a lot of tech from CV
- fixed input size
(although RNNs, in fact, also suffer from it)
- harder to apply to sequence-based tasks
- non-generative

Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org



Read More

Neural Nets for NLP:

<http://cs231n.github.io>

<https://hackernoon.com/the-unreasonable-ineffectiveness-of-deep-learning-in-nlu-e4b4ce3a0da0>

<https://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>

<https://blackboxnlp.github.io/>

Nonlinearities:

<https://towardsdatascience.com/selu-make-fnns-great-again-snn-8d61526802a9>

<https://medium.com/@jaiyamsharma/experiments-with-swish-activation-function-on-mnist-dataset-fc89a8c79ff7>

<https://towardsdatascience.com/the-mostly-complete-chart-of-neural-networks-explained-3fb6f2367464>

<http://building-babylon.net/2017/08/01/hierarchical-softmax/>

Read More x2

Backprop & gradient descent:

<https://medium.com/@karpathy/yes-you-should-understand-backprop-e2f06eab496b>

<http://runder.io/optimizing-gradient-descent/>

<https://openreview.net/pdf?id=ryQu7f-RZ>

<https://fosterelli.co/executing-gradient-descent-on-the-earth>

<https://medium.com/usf-msds/deep-learning-best-practices-1-weight-initialization-14e5c0295b94>

RNN & LSTM:

<https://deeplearning4j.org/lstm.html>

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

<https://medium.com/@aidangomez/let-s-do-this-f9b699de31d9>